

Michael Chang - Senior Machine Learning Engineer

Email: michael.chang@email.com | Phone: +1 415-555-0192 | Location: San Francisco, CA

LinkedIn: linkedin.com/in/michaelchang-ml | GitHub: github.com/mchang-ai

SUMMARY

Senior Machine Learning Engineer with 5+ years experience architecting NLP pipelines, transformer fine-tuning, and scalable PyTorch deployments. Extensive experience in Scikit-learn, Docker, and AWS MLOps.

TECHNICAL SKILLS

Languages & Frameworks: Python, PyTorch, TensorFlow, Scikit-learn, Hugging Face, Transformers, FastAPI

Data & Cloud: PostgreSQL, Redis, Docker, Kubernetes, AWS (S3, EC2, SageMaker), Git, CI/CD

EXPERIENCE

Lead AI Engineer - Vertex Labs (2022 - Present)

Designed real-time semantic search and retrieval-augmented generation (RAG) system processing 10M queries daily using PyTorch, Hugging Face, and Docker. Reduced inference latency by 45%.

ML Engineer - DataCore Systems (2020 - 2022)

Trained classification and anomaly detection models in Python and Scikit-learn on AWS.

EDUCATION

Master of Science in Computer Science - Stanford University (2020)